

STAT 4410/8416 Homework 5

Danuwar Ramesh

Due on Nov 28, 2020

1. **Working with databases:** Please follow the instruction below before answering the questions:

- Install the package `sqldf` using `install.packages('sqldf')`
- Import the library using `library('sqldf')`
- Read the file <https://raw.githubusercontent.com/dsindy/kaggle-titanic/master/data/train.csv> and store it in `titanic`

We can now start writing SQL Script using `SQLDF` library right inside R. See example below:

```
library(sqldf)

sqldf("SELECT passengerid, name, sex
      FROM titanic
      limit 5", drv="SQLite")
```

##	PassengerId	Name	Sex
## 1	1	Braund, Mr. Owen Harris	male
## 2	2	Cummings, Mrs. John Bradley (Florence Briggs Thayer)	female
## 3	3	Heikkinen, Miss. Laina	female
## 4	4	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female
## 5	5	Allen, Mr. William Henry	male

Answer the following questions. Write SQL Script where applicable.

- What is a Database Management System? Name few types of Database Management System. Answer: Database management system is the system that allows to store, modify, and extract information from a database. The different types of database management system are: mysql, oracle, ms sql, db2, ms access etc
- Provide 3 verbs of SQL? Please write what they do Answer: Select: The SELECT statement is used to select data from a database. Update: The SQL UPDATE Query is used to modify the existing records in a table Delete: To remove one or more rows from a table completely, you use the DELETE statement INSERT: used to add new rows of data to a table in the database
- What does the following command do in MySQL?
 - 'show databases; Answer: To display the available databases the command is show.
 - show tables;** Answer: It helps to list the tables for the current/specified database or schema or across the entire account.
- Write SQL script to answer the following questions based on titanic data. Display the results of your script.
 - What is the average age of passengers who survived? Group the data by Sex. Display only the column **Sex, AverageAge**

```
kable(sqldf("SELECT Sex,
              avg(Age) as AverageAge
            FROM titanic
            WHERE Survived = 1
            group by Sex", drv="SQLite"))
```

Sex	AverageAge
female	28.84772
male	27.27602

iv. What is the percentage of passengers who survived in each Passenger Class or `Pclass`? Group the data by `Pclass`.

```
kable(sqldf("SELECT count(Survived) as total_rows_survived
            FROM titanic",
            drv="SQLite"))
```

total_rows_survived
891

```
sqldf("SELECT pclass,
              sex, cast(sum(survived) as real) / count(survived) * 100 as avg_survived_percent
            FROM titanic
            GROUP BY pclass, sex;"
      , drv="SQLite")
```

```
##   Pclass   Sex avg_survived_percent
## 1      1 female          96.80851
## 2      1  male          36.88525
## 3      2 female          92.10526
## 4      2  male          15.74074
## 5      3 female          50.00000
## 6      3  male          13.54467
```

v. What is the average age of all the passenger (survived and not survived)? Group the data by `Pclass`.


```
q1_iii <- sqldf("SELECT Pclass,
                       Sex,
                       Survived,
                       avg(Age) as AverageAge
                  FROM titanic
                  group by Pclass, Sex, Survived", drv="SQLite")

kable(q1_iii)
```

Pclass	Sex	Survived	AverageAge
1	female	0	25.66667
1	female	1	34.93902

Pclass	Sex	Survived	AverageAge
1	male	0	44.58197
1	male	1	36.24800
2	female	0	36.00000
2	female	1	28.08088
2	male	0	33.36905
2	male	1	16.02200
3	female	0	23.81818
3	female	1	19.32979
3	male	0	27.25581
3	male	1	22.27421

```
library(ggplot2)
ggplot(data=q1_iii, aes(x=AverageAge, y=Pclass, color=Survived, group=Survived))+
  geom_line()+
  geom_point()+
  facet_grid(~Sex)+
  xlab("Average Age")+
  ylab("Passenger Class")+
  labs(color='Survived')
```



vi. What is the name, age, sex and pclass of the 5 oldest and 5 youngest persons who died?

```
# 5 Youngest person who died
kable(sqldf(" SELECT Name, Age, Sex, Pclass
  FROM titanic
  WHERE Age IS NOT NULL OR age <= 0
  AND Survived = 0
  ORDER BY Age ASC
  LIMIT 5", drv="SQLite"))
```

Name	Age	Sex	Pclass
Thomas, Master. Assad Alexander	0.42	male	3
Hamalainen, Master. Viljo	0.67	male	2
Baclini, Miss. Helene Barbara	0.75	female	3
Baclini, Miss. Eugenie	0.75	female	3
Caldwell, Master. Alden Gates	0.83	male	2

```
# 5 Oldest person who died
kable(sqldf(" SELECT name, age, sex, pclass
FROM titanic
WHERE age IS NOT NULL OR age <= 0
AND Survived = 0
ORDER BY age DESC
LIMIT 5", drv="SQLite"))
```

Name	Age	Sex	Pclass
Barkworth, Mr. Algernon Henry Wilson	80.0	male	1
Svensson, Mr. Johan	74.0	male	3
Goldschmidt, Mr. George B	71.0	male	1
Artagaveytia, Mr. Ramon	71.0	male	1
Connors, Mr. Patrick	70.5	male	3

vii. On average which Passenger Class is more expensive?

```
kable(sqldf(" SELECT Pclass, avg(Fare) as AverageFare
FROM titanic
GROUP BY Pclass", drv="SQLite"))
```

Pclass	AverageFare
1	84.15469
2	20.66218
3	13.67555

2. **Extracting twitter data;** In this problem we would like to extract data from twitter. For this refer to the documentation in the following link.

<https://github.com/geoffjentry/twitterR/>

- a. **Twitter API set up** Set up twitter API using any of the following methods. Make sure you installed all the packages as mentioned in the class.

Method 1: Read Getting Started section of the above link and create a twitter application by going to the link <https://apps.twitter.com/>. Once you created your application connect twitter from R using the secrets and keys obtained from your twitter application.

```
library(twitterR)
api_key <- "your api key"
```

```
api_secret <- "your api secret"
access_token <- "your access token"
access_token_secret <- "your access token secret"

setup_twitter_oauth(api_key,api_secret,access_token,
                    access_token_secret)
```

Method 2: If you don't like creating an account with twitter and going through all the trouble, you can use my keys (ssh, don't tell anyone). For this download the `hw5-twitter-auth` file from blackboard and load it as follows.

```
load("hw5-twitter-auth")
library(base64enc)
library(httr)
library(twitterR)
setup_twitter_oauth(api_key,api_secret,access_token,
                    access_token_secret)
```

```
## [1] "Using direct authentication"
```

b. Now search twitter messages for “data science job”. Display few job informations.

```
datascience_jobs_tweets <- twitterR::searchTwitter("data science job")
head(datascience_jobs_tweets)
```

```
## [[1]]
## [1] "DJKatie888: RT @TennConserv: ...this was never about science or data; it's about control and plac
##
## [[2]]
## [1] "epuujee: RT @tmj_inh_it: Nervous to apply for a job like \"HCaaS- Data Science- Research and Ser
##
## [[3]]
## [1] "tmj_inh_it: Nervous to apply for a job like \"HCaaS- Data Science- Research and Sensing- Hyderab
##
## [[4]]
## [1] "Sad_Little_King: RT @TennConserv: ...this was never about science or data; it's about control a
##
## [[5]]
## [1] "hbbtruth: RT @TennConserv: ...this was never about science or data; it's about control and plac
##
## [[6]]
## [1] "mrjuoss: RT @freeCodeCamp: Spreadsheet software like Microsoft Excel is used in office work, da
```

c. Search 300 tweets using the hash tag `#chess` and save them in an object called `rTweets`. Show the top 7 sources of tweets (such as android or iphone) in a ordered bar plot.

```
# Search 300 tweets with (#chees)
rTweets <- twitterR::searchTwitter("#chess", n=300)

# Get the source from where the tweet came from
tweet_sources_raw <- sapply(rTweets, function(x) x$statusSource())
```

```

# Remove html tags that is on the sources and only get the value between the tags
# This is the source
library(stringr)
regex_pattern <- '[^>]+\</a>'
tweet_sources_text <- gsub("</a>", "", unlist(str_extract(tweet_sources_raw, regex_pattern)))

# Display few sources
head(tweet_sources_text)

```

```

## [1] "Cheap Bots, Done Quick!" "Instagram"
## [3] "Merchant Media Bot"      "Twitter Web App"
## [5] "TweetDeck"              "The Tweeted Times"

```

```

# We now have the source from which the 300 tweets were being sent
# Now lets get the frequency count of the sources of distict 300 tweets
tweet_sources_text_FreqDist <- table(tweet_sources_text)

# Display few tweets with Frequenct Distribution
head(tweet_sources_text_FreqDist)

```

```

## tweet_sources_text
##      @100DaysOfCode_ 1 2 Chess! Tweetbot      3dRenderBot
##                   1                                10          1
##      AdeptLibrarium      BGA Curator      BirdSite
##                   1                                1          1

```

```

library(ggplot2)
# Convert the Frequency Distirbution data into ta nice Data Frame
# Also sort the frequency data in decreasing order at the same time
library(dplyr)
rTweetsDF_sources <- data.frame(tweet_sources_text_FreqDist) %>% arrange(desc(Freq))

# Display few data
kable(head(rTweetsDF_sources))

```

tweet_sources_text	Freq
Twitter Web App	77
Twitter for iPhone	73
Twitter for Android	56
1 2 Chess! Tweetbot	10
Cheap Bots, Done Quick!	10
Instagram	10

```

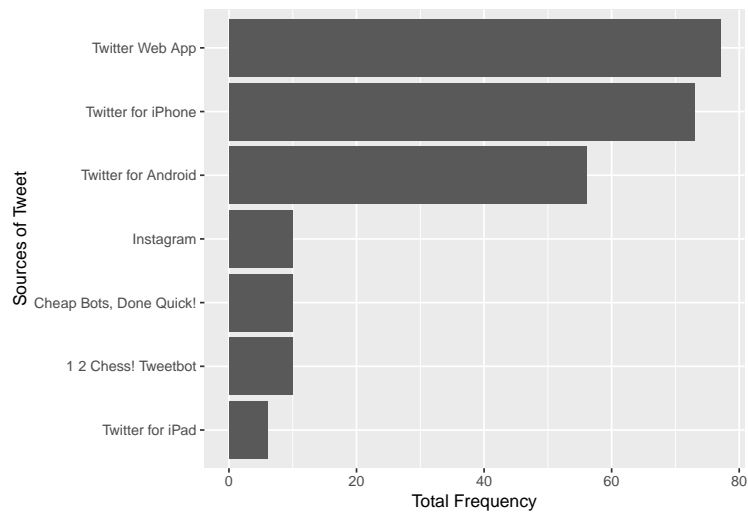
# We only want the top 7 data
# Now lets subset the data
rTweetsDF_sources_top7 <- head(rTweetsDF_sources, 7)

# Display top 7 data
rTweetsDF_sources_top7

```

```
##      tweet_sources_text Freq
## 1      Twitter Web App   77
## 2      Twitter for iPhone 73
## 3      Twitter for Android 56
## 4      1 2 Chess! Tweetbot 10
## 5 Cheap Bots, Done Quick! 10
## 6              Instagram 10
## 7      Twitter for iPad   6
```

```
# Use ggplot to plot the ordered bar plot
ggplot(rTweetsDF_sources_top7, aes(x=reorder(tweet_sources_text, Freq), y =Freq ))+
  geom_bar(stat="identity")+
  xlab("Sources of Tweet")+
  ylab("Total Frequency")+
  coord_flip()
```



- d. Notice that the object `rTweets` is a list. Convert it into a data frame using function `twListToDF` and store it in an object called `dTweets`. Display some data from `dTweets`.

```
dTweets <- twListToDF(rTweets)
head(dTweets)
```

```
##
## 1 Chessbot Results: Black to win in 4 turns\n\u2b1b\u2b1c\u2b1b\u2b1c \u2b1c\u2b1b \n \u2b1b\u2b1c\u2b1b
## 2
## 3
## 4
## 5
## 6
##      favorited favoriteCount replyToSN      created truncated
## 1      FALSE              0      <NA> 2020-11-29 04:29:01      TRUE
## 2      FALSE              1      <NA> 2020-11-29 04:21:23      TRUE
## 3      FALSE              0      <NA> 2020-11-29 04:17:18     FALSE
## 4      FALSE              0      <NA> 2020-11-29 04:15:04     FALSE
## 5      FALSE              0      <NA> 2020-11-29 04:12:39     FALSE
```

```
## 6      FALSE          0      <NA> 2020-11-29 04:07:02      FALSE
##      replyToSID          id replyToUID
## 1      <NA> 1332904375736471554      <NA>
## 2      <NA> 1332902453126238211      <NA>
## 3      <NA> 1332901424594903041      <NA>
## 4      <NA> 1332900862428147712      <NA>
## 5      <NA> 1332900252882386952      <NA>
## 6      <NA> 1332898841532788741      <NA>
##
##                                     statusSource
## 1 <a href="https://cheapbotsdonequick.com" rel="nofollow">Cheap Bots, Done Quick!</a>
## 2      <a href="http://instagram.com" rel="nofollow">Instagram</a>
## 3      <a href="https://merchant.media/home" rel="nofollow">Merchant Media Bot</a>
## 4      <a href="https://mobile.twitter.com" rel="nofollow">Twitter Web App</a>
## 5 <a href="https://about.twitter.com/products/tweetdeck" rel="nofollow">TweetDeck</a>
## 6      <a href="https://www.tweeteditimes.com" rel="nofollow">The Tweeted Times</a>
##      screenName retweetCount isRetweet retweeted longitude latitude
## 1 ChessScenarios          0      FALSE      FALSE      <NA>      <NA>
## 2      GashPhotos          0      FALSE      FALSE     -79.65     43.6
## 3 MerchantMediaCo          3       TRUE      FALSE      <NA>      <NA>
## 4 CodyReedTerry1          0      FALSE      FALSE      <NA>      <NA>
## 5 NewIndianXpress          1       TRUE      FALSE      <NA>      <NA>
## 6 ChessIndiaNet           0      FALSE      FALSE      <NA>      <NA>
```

- e. dTweets has a column showing the time the tweet was created. Generate a plot showing number of tweets on each of the hours. Add a smooth line overlaid on your plot.

```
library(lubridate)

# convert the time into the format ymd_hms()
hour <- data.frame(hour(ymd_hms(dTweets$created)))

# add a new column in the dTweet with hour value
dTweets <- data.frame(hour, dTweets)

# rename the newly added hour column
colnames(dTweets)[1] <- "hourValue"


# filter NA values if any
dTweetsHourCount <- dTweets %>%
  filter(!is.na(hourValue)) %>%
  group_by(hourValue) %>%
  tally()

# Display Tweets per hour table
kable(dTweetsHourCount)
```

hourValue	n
0	11
1	30
2	40
3	28
4	8
18	17

hourValue	n
19	39
20	40
21	41
22	22
23	24

```
# Generate a plot showing number of tweets on each of the hours
# Add a smooth line overlaid on your plot
ggplot(dTweetsHourCount, aes(x=hourValue, y=n))+
  geom_smooth(method = 'loess')+
  geom_point()+
  xlab("Hour Value")+
  ylab("Total tweets per hour")
```



- f. Arrange the dataframe `dTweets` based on the `retweetCount`. While doing this select only columns `text`, `screenName`, `retweetCount`. Store the data in a object called `mostTweets`. Display five texts that are most retweeted.

```
#Arrange the dataframe `dTweets` based on the `retweetCount`
# While doing this select only columns `text, screenName, retweetCount`
mostTweets <- dTweets %>%
  select(text, screenName, retweetCount) %>%
  arrange(desc(retweetCount))

head(mostTweets$text,5)
```

```
## [1] "RT @KhushnumaKashmi: #Chess is the gymnasium of the mind.\nA Chess and carrom competition was o
## [2] "RT @KhushnumaKashmi: Breaking the monotonous routine in view of COVID-19 restrictions, #Indian
## [3] "RT @MikeMavo: unprecedented levels of IQ on display #chess https://t.co/fBalNcdaNf"
## [4] "RT @dgriffinchess: NEWS: My translation of Levenfish's memoir, 'Selected Games & Reminiscen
## [5] "RT @Amb_Salukvadze: The real-life Queen's Gambit: how \nNona Gaprindashvili, world's first fema
```

- g. Generate a bar chart showing top 15 screen names and count of retweets from `mostTweets`. Order the bars based on the retweet counts.

```
top15_mostTweets <- head(mostTweets, 15)

ggplot(top15_mostTweets, aes(x = reorder(screenName, retweetCount), y = retweetCount)) +
  geom_bar(stat= "Identity") +
  xlab("Twitter Screen Name") +
  ylab("Total retweet count") +
  coord_flip() +
  theme(plot.title = element_text(hjust = 0.5))
```

